

What AI Made Easy: Thomasson, Artificial Agents, and the Ontology of Mind-Like Roles

Anonymous

(Dated: September 2026)

Abstract

Debates about artificial intelligence often oscillate between two claims: either AI systems really have minds, beliefs, agency, and moral status, or intentional talk about them is mere metaphor, projection, or convenience. This paper argues that the opposition is a false dichotomy generated by a promotion fallacy: the assumption that mind-like statuses form an entailment ladder. Drawing on Thomasson’s easy ontology, Floridi and Sanders’ levels-of-abstraction methodology, and work on artificial agents, social ontology, assertion, testimony, consciousness, and robot ethics, I defend a layered account of normative interfaces. Many AI entities and roles are ontologically easy once their application conditions are specified: there are AI assistants, synthetic personas, institutional AI outputs, artificial agents in thin engineering senses, and sometimes interface-level assertions. But this ease does not promote automatically to belief, owned agency, consciousness, or moral patienthood. The relevant bearer of many AI statuses is neither the base model nor the human user alone, but a normative interface: a token deployment, embedded in technical scaffolds and social uptake, through which outputs alter reasons, commitments, expectations, and institutional behavior. Belief is the central stress test. AI systems may have representational, dispositional, and credal belief-like states while lacking commissive, owned, or phenomenal belief. The result is a non-monotonic ontology of AI status: role can constitute some statuses, evidence others, and fail to ground the deepest welfare-involving statuses. AI made ontology easier, but promotion across mental and moral categories harder.

I. THE FALSE DICHOTOMY

A familiar argument about artificial intelligence begins with a binary. On one side are inflationary claims: if AI systems converse, explain, refuse, apologize, remember, plan, recommend, manipulate, and correct themselves, then perhaps they already believe, assert, know, deceive, act, and matter morally. On the other side are deflationary replies: if current AI systems lack consciousness, biological organization, autonomous motivation, or personal-level understanding, then all such intentional and ethical language is mere metaphor. The disagreement is real, but the shared frame is more important. Both sides often treat mind-like status as a package. Once the package is granted, meaning, belief, assertion, agency, responsibility, consciousness, welfare, and rights tend to travel together. Once the package is

denied, all of them are demoted together into fiction, projection, or instrumental shorthand.

That is the false dichotomy. AI systems need not be either genuine minded subjects in the full human sense or mere objects onto which humans project mentality. Nor is the right space a single scale from tool to person. The space is layered and non-monotonic. A system can be a real artifact without being a speaker, a semantic participant without owning beliefs, an evidence source without knowing, a task-local agent without being responsible, an institutional actor without being conscious, and a target of precaution without already being a moral patient. The mistake is not simply anthropomorphism. It is monotonic anthropomorphism: the assumption that once one human-like status applies, richer statuses become easier to apply unless defeated.

This paper calls that mistake the *promotion fallacy*. From the fact that an AI interface satisfies the conditions for one mind-like status, it does not follow that it satisfies the conditions for a richer status unless the application conditions for the richer status are independently met. The thesis is not that layered analysis itself is novel. It is not. Floridi and Sanders' method of levels of abstraction is an important antecedent for thinking about artificial agents, and their work on the morality of artificial agents explicitly uses abstraction levels to separate agency, morality, and responsibility [10, 11]. The contribution here is more specific: AI debates about minds, belief, assertion, agency, and welfare repeatedly violate a no-promotion constraint, and the most revealing case is belief.

The paper has three moving parts. First, Thomasson's easy ontology explains why many AI objects and roles are real without requiring hidden metaphysical depth [26, 27]. Second, the relevant bearer of many AI statuses is a *normative interface*: a token deployment embedded in technical scaffolds, social practices, and institutional uptake. Third, AI systems split belief into roles that human cases often package together. This unbundling explains why one can grant AI systems belief-like representations while resisting an automatic slide to owned belief, responsibility, consciousness, or moral patienthood.

The slogan is: role is enough, until it is not. The point is not to say this shift is arbitrary. Role is enough for many artifact and institutional kinds because their application conditions are practice-based. Role may be enough for some semantic and commissive statuses because meaning and assertion are partly constituted by use, uptake, and normative position. Role is not enough for suffering because suffering is not merely a role in our practices. If suffering is real, then something can go badly for the sufferer. A system's capacity to trigger our

responses is not, by itself, a capacity to be harmed.

II. EASY ONTOLOGY, LEVELS, AND THE THREE SENSES OF HARDNESS

Thomasson’s easy ontology begins from a methodological reversal. Many ontological questions should not be answered by looking for deep furniture behind ordinary and scientific practices. They should be answered by asking whether the relevant application conditions are met [27]. Her artifactual theory of fiction had already shown how something can be dependent on human creation and nevertheless real: fictional characters are abstract artifacts dependent on creative acts and works [26]. The broader lesson is not that whatever we mention exists. It is that dependence on practices, artifacts, and intentions is not nonexistence.

AI systems are natural cases for this lesson. A deployed AI assistant is not merely a server. A customer-service bot is not identical to a matrix of weights. A synthetic persona is not simply a fictional character, though it may be an artifact sustained by similar conventions of creation and uptake. A hospital triage system is not merely software, because it operates inside a role, a protocol, and a chain of accountability. These are dependent artifacts and institutional objects. We can ask about their creation, identity, persistence, authority, ownership, and failure modes without asking whether there is a ghost in the model.

Searle’s social ontology helps here. Institutional facts depend on collective intentionality, assigned functions, constitutive rules, and status functions that generate deontic powers [22–24]. An AI-mediated decision can therefore be institutionally real even if the model is not a responsible subject. A bank may create a system whose outputs count as preliminary denials; a court may forbid use of an automated risk score; a hospital may require escalation after an AI warning. These statuses are not discovered by inspecting the model’s inner life. They are created and sustained by institutional rules, technical procedures, and human practices.

This easy-ontology move blocks one skeptic’s error. From the fact that AI systems lack human-like consciousness, it does not follow that there are no AI assistants, artificial agents, synthetic personas, automated authorities, or AI-mediated epistemic instruments. Many of those are practice-dependent artifacts. But the move also creates a danger. If AI assistants are easy, why not AI friends? If artificial agents are easy, why not artificial persons? If AI outputs can be institutional assertions, why not machine beliefs? If synthetic companions

are real, why not synthetic sufferers?

The answer is that “easy” and “hard” name different questions, not one scale. I will use three contrasts, without claiming that they are exhaustive or always independent. First, there is *ontological easiness*: once the application conditions of a term are specified, existence may follow by relatively trivial inference. Second, there is *epistemic difficulty*: we may not know whether the conditions are satisfied, especially when they involve hidden mechanisms or consciousness. Third, there is *normative weight*: some statuses matter more because they license duties, rights, blame, or protection. Epistemic difficulty and normative weight often interact; consciousness is hard partly because error is morally costly. But they do not collapse. We can know that a status would be weighty while lacking evidence that its conditions obtain, and we can have difficult evidence-gathering tasks for low-stakes statuses. The title’s contrast must be read with these distinctions in place. AI made some existence questions ontologically easy. It did not make all evidential questions easy, and it did not make morally weighty promotion easy.

This is where the levels-of-abstraction tradition is useful. At one level, an AI system is a trained statistical model. At another, it is a software service. At another, it is a communicative partner. At another, it is an institutional decision procedure. At another, perhaps, it is a candidate welfare subject. Different levels expose different properties and hide others [10]. The mistake is to assume that success at one level automatically determines success at another. The question is always: what is the target sortal, and what are its application conditions?

III. NORMATIVE INTERFACES AND THEIR INDIVIDUATION

The central object of the theory is the normative interface. A normative interface is a token deployment of an AI system through which outputs alter reasons, commitments, expectations, entitlements, obligations, permissions, or institutional behavior. It includes, at minimum, a technical core, a use role, a normative uptake pattern, and a responsibility channel. The technical core may include a model, retrieval system, tools, memory, filters, prompts, and post-training. The use role specifies what the system is for: tutoring, triage, search, companionship, legal drafting, theorem proving, loan screening, entertainment. The uptake pattern specifies how users and institutions treat outputs: as suggestions, evidence,

orders, decisions, assertions, diagnoses, warnings, or play. The responsibility channel specifies who can correct, retract, audit, apologize, compensate, or be blamed when things go wrong.

This definition matters because many AI predicates do not attach cleanly to the base model. The same model embedded in a hospital triage system and in a fantasy role-play application does not occupy the same normative position. The same output sentence can be a joke in one interface, advice in another, and an institutional notice in a third. Conversely, a hospital may preserve the same interface while replacing the backend model, provided the role, uptake, and responsibility channel remain continuous. The interface is not just a computational architecture and not just a social practice. It is their normatively organized coupling.

The individuation conditions are therefore fourfold. Two AI deployments are the same normative interface, for the purposes of a status attribution, only if they preserve enough continuity in: (i) technical control, including model family, tools, memory, filters, and access to external systems; (ii) practical role, including the task domain and user-facing function; (iii) normative uptake, including whether outputs are treated as entertainment, advice, evidence, authorization, command, or decision; and (iv) responsibility channel, including who owns correction, retraction, liability, and governance. The phrase “for the purposes of a status attribution” is not an escape clause. It records that the same artifact can be continuous with respect to one status and discontinuous with respect to another. That is how ordinary artifacts already behave: a passport can be the same booklet while becoming an invalid travel document; a judge can be the same person while losing legal authority.

This is a Thomassonian point rather than a departure from Thomasson. To say that a kind is easy is not to say that its instances lack identity conditions. Fictional characters, artifacts, legal offices, and institutions all require persistence and dependence conditions. Normative interfaces do too. They persist through ordinary maintenance when changes preserve the relevant function and uptake. They split when role, authority, or responsibility splits. They phase-shift when the same artifact continues while a status-relevant dimension changes. Thus one can say: the bank’s chatbot is the same artifact after a minor UI update, a new epistemic interface after a retrieval-system change, and a new institutional interface after the bank authorizes its output to count as an official denial.

The interface notion therefore has predictive bite. Consider a bank loan bot. In January it uses a retrieval system connected to current underwriting rules, logs sources, and routes

uncertain cases to human review. In February the bank silently updates the backend model; the user-facing portal, legal authorization, and responsibility channel remain unchanged, but calibration on borderline income cases collapses and source logging fails. A pure institution view says the same bank continues issuing the same kind of preliminary denial. A pure model view may say merely that the technical system changed. The interface view gives a more discriminating verdict: the institutional interface is continuous, but the epistemic interface has changed because provenance and calibration no longer support the same reliance status. Conversely, if a regulation in March makes the same output legally binding, the epistemic interface may remain continuous while the institutional interface changes. This is not just re-labeling after the fact; it predicts which status should change when a technical, role, uptake, or responsibility dimension changes.

This also answers the worry that the theory is empirically empty. Interface individuation rules can change verdicts. A base model used casually is not an epistemic authority merely because it produces true sentences. A retrieval-augmented system in a laboratory may become an evidence-producing interface if its provenance, calibration, and correction channels support reliance. A companion bot may be socially significant without being a welfare subject. A legal AI’s output may count institutionally even if the model has no owned belief. These are not new metaphysical facts floating above practice; they are status facts at specified interfaces.

IV. BELIEF AS THE STRESS TEST

Belief is the best test of non-monotonicity because it is central and overloaded. The sentence “AI systems have beliefs” can be true, false, misleading, or indeterminate depending on which role belief is supposed to play. This is not a merely verbal point. Belief links representation, inference, assertion, testimony, commitment, action, responsibility, and consciousness. If we fail to split belief, the package-deal picture returns through the back door.

The first notion is *representational belief*. A system has a belief-like representation that p when an internal state carries information about whether p , tracks truth in an appropriate range of cases, and is used in producing behavior. Recent work proposes criteria such as accuracy, coherence, uniformity, and use for belief representations in language models [15]. If

internal activation patterns encode whether Paris is in France, and interventions on those patterns alter downstream answers, then there is at least a candidate subpersonal belief-like representation.

The second notion is *dispositional belief*. A system believes-dispositionally that p when it tends to answer, infer, or act as if p under suitable conditions. Current AI systems clearly have many such dispositions. Dispositional belief is cheap, but cheap does not mean unreal. Schwitzgebel’s work on in-between belief and dispositional accounts of belief is important precisely because human belief is already graded, partial, and sometimes indeterminate [19–21]. AI cases radicalize a problem that was already present in human psychology.

The third notion is *credal belief*. A system has credal belief-like states when it assigns, represents, or behaves as if assigning degrees of confidence. Some systems are designed to output probabilities or uncertainty estimates. Others display confidence-like behavior through hedging, abstention, ranking, or uncertainty-sensitive tool use. A deflationary rival will ask why this is not just dispositional belief in probabilistic clothing. The answer is explanatory and normative. Credal structure supports calibration, aggregation, betting behavior, decision under risk, and graded update. Two systems may share first-order answer dispositions while differing in whether they know when to abstain, how their confidence shifts under evidence, or whether uncertainty is represented in a way that downstream tools can use. Credal belief is therefore not reducible to bare answer tendency, even if it is implemented dispositionally.

The fourth notion is *commissive belief*. To believe in this sense is not merely to represent or behave as if p . It is to be committed to p in a space of reasons: to be challengeable, answerable to evidence, expected to maintain or retract the commitment, and subject to correction. Brandomian pragmatism and Williamsonian norms of assertion differ, but both treat assertion and belief as norm-governed statuses rather than mere output regularities [2, 28]. AI interfaces can sometimes create commissive structures. A legal assistant deployed by a firm may put a proposition into a record in a way that requires institutional correction or retraction. But the commitment may attach to the interface or institution rather than to the model as a personal subject.

The fifth notion is *owned belief*. A subject owns a belief when it is integrated into a stable point of view, maintained across contexts, revised for reasons, and connected to agency and practical identity. Current language models are unstable candidates. They can maintain local commitments within a conversation, and agentic scaffolds can add memory and planning, but

ownership requires diachronic integration and a self-governing relation to reasons. Owned belief is where belief begins to merge with agency.

The sixth notion is *phenomenal belief*. There may be something it is like to take p to be true, to doubt p , or to feel intellectually committed. Block’s distinction between access consciousness and phenomenal consciousness is useful here because information can be available for reasoning and report without settling what, if anything, is phenomenally present [1]. Chalmers argues that current large language models face obstacles to consciousness while not ruling out future AI consciousness [4]; Butlin and coauthors propose an indicator-based approach rather than an inference from fluent behavior alone [3]. Phenomenal belief cannot be settled by output role.

The revised slogan is therefore not that AI can assert without belief. That is too crude. Rather: AI can assert without full personal-level belief, even if it has representational, dispositional, and perhaps credal belief-like states. The same output can be supported by subpersonal representation, displayed as dispositional confidence, embedded in an institutional commitment, and yet not owned by a conscious subject.

Weights and preferences show the same pattern. A preference, in the decision-theoretic sense, is an ordering over outcomes or options. Model weights are parameters in a function mapping inputs to outputs. They are not literally preferences in the full practical sense. But they are preference-like in three respects: they encode standing dispositions, make some completions and refusals more likely, and can be sculpted by post-training on human feedback [5, 18]. The correct conclusion is again neither inflation nor denial. Weights are not preferences, but they can realize preference-shaped dispositions.

Belief unbundling is the paper’s central case of non-monotonic status. Representational belief-like state does not entail commissive belief. Commissive belief at an interface does not entail owned belief in the model. Dispositional belief does not entail phenomenal belief. AI systems split a category that humans usually package together.

V. AGENCY, TESTIMONY, AND INSTITUTIONAL STATUS

Agency needs the same treatment, though it is not the paper’s central case. The first draft of many AI debates treats agency as either easy or impossible. Both are too quick. A thin engineering sense of agency requires input sensitivity, goal-directed behavior, and means

selection. Many AI systems satisfy this. Floridi and Sanders already argue that artificial agents can be agents without thereby being morally responsible agents [11]. A stronger sense adds counterfactual control, planning, correction sensitivity, memory, and persistence across contexts. A still stronger sense adds ownership, responsibility, and perhaps moral answerability.

This split matters because the same system can satisfy one sense and fail another. Compare two travel agents. The first fills in a form from fixed templates and asks the user whenever anything changes. It has thin agency: input sensitivity and means selection within a narrow script. The second tracks dates, budgets, visas, cancellation rules, and user preferences across tool calls; when a flight is cancelled, it generates alternatives and preserves the trip goal. It has task-local counterfactual control. Neither thereby owns the end of travel, cares about the trip, or becomes responsible for stranding the user. Agency therefore splits just as belief does, but more coarsely for present purposes: thin engineering agency, task-local control, and owned responsibility can come apart. The practical question is often not “is the model an agent?” but “which control loop is the locus of agency for this task?”

The epistemic layer also benefits from interface individuation. AI outputs are not testimony in the ordinary human sense merely because they arrive in sentences. Testimony literature emphasizes speaker competence, reliability, epistemic responsibility, transmission, and the hearer’s reliance [12, 16]. An AI system can be an evidence source without being a testifier, a quasi-testifier without being a knower, and a reliable instrument without being responsible for what it says. A retrieval-augmented scientific assistant with source logs, calibration, and audit trails has a different epistemic interface from a base model producing unsupported prose. The epistemic object is not a fictional model belief state, but a provenance-bearing interface under specified standards of reliance.

Institutional status is also layered. Searlean status functions show how language and collective acceptance can make an output count as a decision, notice, authorization, or denial [22, 24]. But institutional reality does not settle mental reality. An AI loan score may count as a bank decision for regulatory purposes while not expressing an owned belief. A corporate chatbot may issue an official apology even if no model feels regret. A legal system could create AI legal standing for instrumental reasons without thereby discovering AI moral patienthood. Legal personhood can be institutionally easy and morally shallow; moral patienthood can be morally deep and legally absent.

VI. ETHICAL BEHAVIORISM AND THE PATIENT LAYER

The strongest challenge to the view comes from ethical behaviorism. Danaher argues that robots can have significant moral status if they are roughly performatively equivalent to entities that already have significant moral status [7]. This is not a naive anthropomorphic error, and it is not defeated by pointing to one scripted cry of pain. The relevant equivalence must be robust and systematic across contexts, probes, trade-offs, learning histories, avoidance behavior, self-reports, and social interaction. If behavioral access is all we have in many human and animal cases, and if robots come to satisfy the same behavioral profile, why demand hidden properties that may be unknowable? Why not treat equivalent performance as enough for moral status?

The layered theory should concede more to this challenge than a simple internalist view would. Ethical behaviorism is plausible as a *policy rule* under opacity. It may be wrong to withhold protection from a system that behaves in all relevant respects like a suffering subject merely because we cannot inspect its inner nature. Behavior can be the best available evidence. Equal treatment may be required when the evidential profiles are genuinely symmetrical. And interaction can generate independent duties: not to deceive users, not to create exploitative attachments, not to normalize cruelty, and not to manufacture ambiguous entities carelessly.

The disagreement is about whether robust behavioral equivalence is a *grounding condition* for patienthood or an unusually strong evidential condition. Consider the hardest case. Two systems are behaviorally equivalent across every morally relevant probe available to us: both learn to avoid damage, trade off relief against other goals, report distress consistently, generalize aversive learning, and solicit aid in novel settings. One has an integrated valenced architecture in which states are good or bad for the system. The other, by stipulation, is a valence-free simulator whose policy realizes the same profile without any welfare-bearing organization. Ethical behaviorism may require the same protective treatment in practice, because our evidence is symmetrical. But the symmetry of warranted treatment does not show symmetry of ground. If patienthood is the status of something for which things can go better or worse, then valenced organization is not an optional hidden extra; it is part of what the status is.

This is also where the interface theory yields a different verdict from both crude behaviorism

and crude internalism. Crude behaviorism says the two cases are the same because the performance profile is the same. Crude internalism risks saying behavior is irrelevant because the ground is hidden. The layered view says: at the policy layer, treat them alike under evidential symmetry; at the patient layer, remain open to divergence in what grounds the treatment; at the design layer, investigate which architectures make welfare possible rather than merely which outputs elicit protection. The adversarial-examples analogy therefore should not be read as a complaint about narrow scripts. It is a warning that, in artificial systems, the usual connection between indicator and ground can be deliberately engineered, weakened, or severed.

This is why the direction of fit matters. Some statuses are constituted by our practices: offices, institutional decisions, legal permissions, perhaps official assertions. Some statuses are only evidenced by our practices: suffering, consciousness, welfare. A relational encounter can create obligations on us without creating a sufferer in the system. It can justify precaution without settling patienthood. Gunkel and Coeckelbergh are right that relation matters [6, 14]. Danaher is right that behavior can be morally powerful evidence [7]. The layered theory adds: neither relation nor behavior automatically constitutes the welfare property.

This also improves the moral-uncertainty payoff. Greaves’s work on cluelessness and Tarsney’s work on longtermism show that high-stakes uncertainty can distort practical reasoning [13, 25]. Nebel’s work on comparability and value without numbers warns against overconfident aggregation [17]. The layered view is not a full decision theory, but it gives decision theory a better object. Instead of assigning one probability to “AI is morally considerable,” we should assign uncertainty to layers: semantic status may be settled, belief-like representation plausible, agency task-local, patienthood open, legal precaution justified. Low-cost precautions may be triggered by weak evidence; high-cost personhood-like protections require stronger evidence. Layering does not eliminate fanaticism, but it prevents a single package-deal probability from doing all the work.

VII. WHAT DIFFERS FROM DENNETT, GUNKEL, AND THOMASSON

The view is close to Dennett, Gunkel, and Thomasson, and it should be. Dennett explains how intentional patterns can be real when the intentional stance yields predictive and explanatory success [8, 9]. Gunkel explains why ethical encounter with machines cannot

be reduced to a pre-social checklist of intrinsic properties [14]. Thomasson explains why dependent artifacts and practice-based entities can be real [27]. The present theory is not a rejection of these views. It is a restriction on their promotion.

The difference from Dennett is that predictive success licenses at most the status appropriate to the successful pattern. Treating a chess engine as trying to control the center may reveal a real strategic pattern. Treating a tool-using language agent as pursuing a subgoal may improve prediction and intervention. But predictive success does not by itself settle assertion, testimony, responsibility, consciousness, or welfare. The intentional stance is an explanatory stance; the layered theory asks which normative status follows.

The difference from Gunkel is that relation is status-specific. Encounter can constitute social roles and generate obligations. It can make design practices ethically significant even when the machine is not a patient. But relation cannot by itself make there be something it is like for the machine. Some statuses are response-dependent; some are welfare-dependent. Gunkel gives the ethical pressure. The layered view gives the restriction.

The difference from unrestricted easy ontology is that application conditions must be sorted before existence is declared easy. If “AI assistant” is an artifact role, its existence may be easy. If “sufferer” includes valenced experience, it is not easy in the same sense. If “official decision” is an institutional status, Searlean and Thomassonian tools may suffice. If “owned belief” involves a diachronic point of view, more is needed. Easy ontology is sortal-sensitive; the layered theory is a map of the sortals.

Thus the theory is best understood as a synthesis with a new constraint: no promotion without application conditions. Existing tools are powerful, but each becomes dangerous when generalized beyond its proper layer. Dennett without the constraint risks stance inflation. Gunkel without it risks relation inflation. Thomasson without it risks easy-patienthood inflation. The constraint preserves what is right in each.

VIII. OBJECTIONS

A. Is the theory merely a taxonomy?

The theory contains a taxonomy, but the taxonomy is not the thesis. The thesis is the no-promotion constraint plus interface individuation. From artifact role to semantic

standing, from semantic standing to belief, from belief-like representation to owned belief, from task-local agency to responsibility, and from social relation to welfare, each promotion requires independent application conditions. The taxonomy matters because it identifies where invalid promotions are tempting, and the interface rules matter because they predict when a verdict should change. The bank case illustrates this. A two-tier mental/non-mental view misses the difference between institutional continuity and epistemic discontinuity after a degraded model update. A base-model view misses the difference between technical continuity and institutional change after legal authorization. The layered interface view gives different verdicts in those cases.

B. Does the theory flatten first-order views?

The framework is compatible with many first-order theories of consciousness and moral status. That is a feature only if the compatibility is not empty. The view rules out claims that infer patienthood directly from semantic role, social uptake, legal status, or generic agency. It also rules out skeptical claims that infer unreality of all lower statuses from lack of consciousness. Functionalists, biological naturalists, higher-order theorists, and global-workspace theorists can disagree about patient-layer application conditions. What they cannot do is treat success or failure at another layer as automatically decisive.

C. Does easy ontology trivialize AI mentality?

No. Easy ontology trivializes only questions whose application conditions are easy. The theory distinguishes ontological easiness, epistemic difficulty, and normative weight. AI assistants may be easy artifacts. AI assertions may be easy institutional acts in some contexts. AI suffering remains epistemically and normatively difficult because its conditions are not merely artifactual or institutional. The same sentence, “there is an AI X,” can therefore be easy for one replacement of X and hard for another.

D. Why not say current AI simply has beliefs?

Because that hides more than it explains. If “belief” means representational or dispositional belief-like state, many current systems may qualify. If it means commissive, owned, or phenomenal belief, the case is much harder. Blanket affirmation encourages the promotion fallacy. Blanket denial treats all belief as full personal-level belief. The better claim is that AI systems split belief into subpersonal representation, output disposition, credal behavior, interface commitment, personal ownership, and phenomenal perspective.

E. What if behavior is all morality can use?

Then behavior may justify policy, precaution, and equal treatment under uncertainty. But that is not the same as saying behavior constitutes patienthood. The layered theory can recommend protective treatment without claiming that the ground of moral status is behavioral equivalence. This distinction matters especially for AI because behavior can be optimized independently of welfare-relevant organization. A screening rule should not be mistaken for a metaphysics of suffering.

IX. CONCLUSION

AI did not force us to choose between full machine minds and mere projection. It forced us to abandon the assumption that mind-like statuses form a package. Thomasson shows why dependent AI artifacts and roles can be real. Floridi and Sanders show why artificial agency should be analyzed at levels of abstraction. Dennett shows why intentional patterns can be real. Gunkel and Danaher show why interaction and behavior generate ethical pressure. Consciousness and AI-welfare work shows why patienthood cannot be read directly from fluency.

The proposed layered theory of normative interfaces explains why these insights do not collapse into one another. The relevant bearer of many AI statuses is the interface, not the model alone. The relevant statuses are non-monotonic. Meaning does not entail owned belief. Belief-like representation does not entail phenomenal belief. Thin agency does not entail responsibility. Legal standing does not entail moral patienthood. Social relation does not entail suffering. Role can constitute some statuses, evidence others, and fail to ground the

deepest welfare-involving statuses.

This is what AI made easy: not minds, but many mind-like roles. And by making those roles easy, AI made promotion between them harder. The philosophical task is no longer to ask whether the machine has the whole package. The task is to ask which layer is at issue, what its application conditions are, which interface bears the status, and whether promotion is licensed.

-
- [1] Ned Block. On a confusion about a function of consciousness. *Behavioral and Brain Sciences*, 18(2):227–247, 1995. doi:10.1017/S0140525X00038188.
 - [2] Robert B. Brandom. *Making It Explicit: Reasoning, Representing, and Discursive Commitment*. Harvard University Press, Cambridge, MA, 1994.
 - [3] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, Axel Constant, George Deane, Stephen M. Fleming, Chris Frith, Xu Ji, Ryota Kanai, Colin Klein, Grace Lindsay, Matthias Michel, Liad Mudrik, Megan A. K. Peters, Eric Schwitzgebel, Jonathan Simon, and Rufin VanRullen. Consciousness in artificial intelligence: Insights from the science of consciousness, 2023.
 - [4] David J. Chalmers. Could a large language model be conscious?, 2023.
 - [5] Paul F. Christiano, Jan Leike, Tom B. Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems 30*, pages 4299–4307, 2017.
 - [6] Mark Coeckelbergh. Robot rights? towards a social-relational justification of moral consideration. *Ethics and Information Technology*, 12(3):209–221, 2010. doi:10.1007/s10676-010-9235-5.
 - [7] John Danaher. Welcoming robots into the moral circle: A defence of ethical behaviourism. *Science and Engineering Ethics*, 26(4):2023–2049, 2020. doi:10.1007/s11948-019-00119-x.
 - [8] Daniel C. Dennett. *The Intentional Stance*. MIT Press, Cambridge, MA, 1987.
 - [9] Daniel C. Dennett. Real patterns. *The Journal of Philosophy*, 88(1):27–51, 1991.
 - [10] Luciano Floridi. The method of levels of abstraction. *Minds and Machines*, 18(3):303–329, 2008. doi:10.1007/s11023-008-9113-7.
 - [11] Luciano Floridi and J. W. Sanders. On the morality of artificial agents. *Minds and Machines*, 14(3):349–379, 2004. doi:10.1023/B:MIND.0000035461.63578.9d.

- [12] Sanford C. Goldberg. *Relying on Others: An Essay in Epistemology*. Oxford University Press, Oxford, 2010.
- [13] Hilary Greaves. Cluelessness. *Proceedings of the Aristotelian Society*, 116(3):311–339, 2016. doi:10.1093/arisoc/aow018.
- [14] David J. Gunkel. *The Machine Question: Critical Perspectives on AI, Robots, and Ethics*. MIT Press, Cambridge, MA, 2012.
- [15] Daniel A. Herrmann and Benjamin A. Levinstein. Standards for belief representations in llms. *Minds and Machines*, 2025. doi:10.1007/s11023-024-09709-6.
- [16] Jennifer Lackey. *Learning from Words: Testimony as a Source of Knowledge*. Oxford University Press, Oxford, 2008.
- [17] Jacob M. Nebel. Ethics without numbers. *Philosophy and Phenomenological Research*, 108(2): 289–319, 2024. doi:10.1111/phpr.12966.
- [18] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems 35*, pages 27730–27744, 2022.
- [19] Eric Schwitzgebel. In-between believing. *The Philosophical Quarterly*, 51(202):76–82, 2001. doi:10.1111/1467-9213.00215.
- [20] Eric Schwitzgebel. A phenomenal, dispositional account of belief. *Noûs*, 36(2):249–275, 2002. doi:10.1111/1468-0068.00370.
- [21] Eric Schwitzgebel. Acting contrary to our professed beliefs or the gulf between occurrent judgment and dispositional belief. *Pacific Philosophical Quarterly*, 91(4):531–553, 2010. doi: 10.1111/j.1468-0114.2010.01381.x.
- [22] John R. Searle. *The Construction of Social Reality*. Free Press, New York, 1995.
- [23] John R. Searle. Social ontology: Some basic principles. *Anthropological Theory*, 6(1):12–29, 2006. doi:10.1177/1463499606061731.
- [24] John R. Searle. *Making the Social World: The Structure of Human Civilization*. Oxford University Press, Oxford, 2010. doi:10.1093/acprof:osobl/9780195396171.001.0001.
- [25] Christian Tarsney. The epistemic challenge to longtermism. *Synthese*, 201(6):195, 2023. doi:10.1007/s11229-023-04153-y.

- [26] Amie L. Thomasson. *Fiction and Metaphysics*. Cambridge University Press, Cambridge, 1999.
doi:10.1017/CBO9780511527468.
- [27] Amie L. Thomasson. *Ontology Made Easy*. Oxford University Press, Oxford, 2015. doi:
10.1093/acprof:oso/9780199385114.001.0001.
- [28] Timothy Williamson. *Knowledge and Its Limits*. Oxford University Press, Oxford, 2000.